

## Wine: Chemical Analysis

### Full-Out Linear Regression

- Model
- Diagnostics
- Transformation
- Random Effect Model

## I

I fit a linear model with quality as a response and the rest of variables as predictors, but I noticed that response is a discrete variable which can be thought of as being treated a continuous variable in this regression, and the problem naturally arises later in residual vs fit plot.

```
> summary(lmod)
```

Call:

```
lm(formula = quality ~ facidity + vacity + citric + rsugar +
    chlorides + fso2 + tso2 + density + pH + so4 + alcohol)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.68911	-0.36652	-0.04699	0.45202	2.02498

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	2.197e+01	2.119e+01	1.036	0.3002
facity	2.499e-02	2.595e-02	0.963	0.3357
vacity	-1.084e+00	1.211e-01	-8.948	< 2e-16 ***
citric	-1.826e-01	1.472e-01	-1.240	0.2150
rsugar	1.633e-02	1.500e-02	1.089	0.2765
chlorides	-1.874e+00	4.193e-01	-4.470	8.37e-06 ***
fso2	4.361e-03	2.171e-03	2.009	0.0447 *
tso2	-3.265e-03	7.287e-04	-4.480	8.00e-06 ***
density	-1.788e+01	2.163e+01	-0.827	0.4086
pH	-4.137e-01	1.916e-01	-2.159	0.0310 *
so4	9.163e-01	1.143e-01	8.014	2.13e-15 ***
alcohol	2.762e-01	2.648e-02	10.429	< 2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.648 on 1587 degrees of freedom

Multiple R-squared: 0.3606, Adjusted R-squared: 0.3561

F-statistic: 81.35 on 11 and 1587 DF, p-value: < 2.2e-16

I perform model selection by backward selection based on elimination of variable with the largest p-value.

The largest p-value is with the variate density

```
> lmod1 = update(lmod, .~., -density)
> summary(lmod1)
```

Call:

```
lm(formula = quality ~ facidity + vacidity + citric + rsugar +
    chlorides + fso2 + tso2 + pH + so4 + alcohol)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.67204	-0.36527	-0.04523	0.45628	2.03894

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	4.4538341	0.6125783	7.271	5.59e-13	***
facidity	0.0081441	0.0160586	0.507	0.61212	
vacidity	-1.0964449	0.1200866	-9.130	< 2e-16	***
citric	-0.1836098	0.1471561	-1.248	0.21232	
rsugar	0.0089507	0.0120542	0.743	0.45787	
chlorides	-1.9067341	0.4173928	-4.568	5.30e-06	***
fso2	0.0045147	0.0021631	2.087	0.03704	*
tso2	-0.0033120	0.0007264	-4.560	5.52e-06	***
pH	-0.5042762	0.1571117	-3.210	0.00136	**
so4	0.8928974	0.1107548	8.062	1.46e-15	***
alcohol	0.2927427	0.0173394	16.883	< 2e-16	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6479 on 1588 degrees of freedom

Multiple R-squared: 0.3603, Adjusted R-squared: 0.3562

F-statistic: 89.43 on 10 and 1588 DF, p-value: < 2.2e-16

The largest p-value is now with the variate rsugar

```
> lmod2 <- lm(quality ~ facidity + vacidity + citric + chlorides +
fso2 + tso2 + pH + so4 + alcohol)
> summary(lmod2)
```

Call:

```
lm(formula = quality ~ facidity + vacidity + citric + chlorides +
    fso2 + tso2 + pH + so4 + alcohol)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.68601	-0.36723	-0.04516	0.45629	2.02723

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	4.4410923	0.6122514	7.254	6.31e-13	***
facidity	0.0090201	0.0160129	0.563	0.57331	
vacidity	-1.0905804	0.1198096	-9.103	< 2e-16	***
citric	-0.1756500	0.1467444	-1.197	0.23149	
chlorides	-1.8893071	0.4166737	-4.534	6.21e-06	***
fso2	0.0046664	0.0021532	2.167	0.03036	*
tso2	-0.0032560	0.0007224	-4.507	7.04e-06	***
pH	-0.5022333	0.1570654	-3.198	0.00141	**
so4	0.8872849	0.1104810	8.031	1.86e-15	***
alcohol	0.2940206	0.0172514	17.043	< 2e-16	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6479 on 1589 degrees of freedom

Multiple R-squared: 0.3601, Adjusted R-squared: 0.3564

F-statistic: 99.34 on 9 and 1589 DF, p-value: < 2.2e-16

```
> lmod3 <- lm(quality ~ vacidity + citric + chlorides + fso2 + tso2 +
pH + so4 + alcohol)
> summary(lmod3)
```

Call:

```
lm(formula = quality ~ vacidity + citric + chlorides + fso2 +
    tso2 + pH + so4 + alcohol)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.66890	-0.37044	-0.04474	0.45697	2.02363

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	4.6680876	0.4608410	10.129	< 2e-16	***
vacidity	-1.0736123	0.1159362	-9.260	< 2e-16	***
citric	-0.1295444	0.1217717	-1.064	0.2876	
chlorides	-1.9494185	0.4026906	-4.841	1.42e-06	***
fso2	0.0047601	0.0021463	2.218	0.0267	*
tso2	-0.0033658	0.0006954	-4.840	1.42e-06	***
pH	-0.5491501	0.1331350	-4.125	3.90e-05	***

```

so4          0.8914283  0.1102122   8.088 1.19e-15 ***
alcohol      0.2928780  0.0171280  17.099 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1590 degrees of freedom
Multiple R-squared:  0.3599,    Adjusted R-squared:  0.3567
F-statistic: 111.8 on 8 and 1590 DF,  p-value: < 2.2e-16

> lmod4 <- lm(quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4
+ alcohol)
> summary(lmod4)

Call:
lm(formula = quality ~ vacidity + chlorides + fso2 + tso2 + pH +
    so4 + alcohol)

Residuals:
    Min       1Q   Median       3Q      Max
-2.68918 -0.36757 -0.04653  0.46081  2.02954

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  4.4300987   0.4029168   10.995 < 2e-16 ***
vacidity     -1.0127527   0.1008429  -10.043 < 2e-16 ***
chlorides    -2.0178138   0.3975417   -5.076 4.31e-07 ***
fso2          0.0050774   0.0021255    2.389  0.017 *
tso2         -0.0034822   0.0006868   -5.070 4.43e-07 ***
pH           -0.4826614   0.1175581   -4.106 4.23e-05 ***
so4           0.8826651   0.1099084    8.031 1.86e-15 ***
alcohol       0.2893028   0.0167958   17.225 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1591 degrees of freedom
Multiple R-squared:  0.3595,    Adjusted R-squared:  0.3567
F-statistic: 127.6 on 7 and 1591 DF,  p-value: < 2.2e-16

```

We have eliminated the following variables in order: density, residual sugar, fixed acidity, and citric acid.

No other variables have high enough p-value to be eliminated anymore.

Using backward elimination based on AIC value also reaches the same model.

```
> step(lmod)
```

```
Start: AIC=-1375.49
```

```
quality ~ facidity + vacidity + citric + rsugar + chlorides +  
fso2 + tso2 + density + pH + so4 + alcohol
```

	Df	Sum of Sq	RSS	AIC
- density	1	0.287	666.70	-1376.8
- facidity	1	0.389	666.80	-1376.5
- rsugar	1	0.498	666.91	-1376.3
- citric	1	0.646	667.06	-1375.9
<none>			666.41	-1375.5
- fso2	1	1.694	668.10	-1373.4
- pH	1	1.957	668.37	-1372.8
- chlorides	1	8.391	674.80	-1357.5
- tso2	1	8.427	674.84	-1357.4
- so4	1	26.971	693.38	-1314.0
- vacidity	1	33.620	700.03	-1298.8
- alcohol	1	45.672	712.08	-1271.5

```
Step: AIC=-1376.8
```

```
quality ~ facidity + vacidity + citric + rsugar + chlorides +  
fso2 + tso2 + pH + so4 + alcohol
```

	Df	Sum of Sq	RSS	AIC
- facidity	1	0.108	666.81	-1378.5
- rsugar	1	0.231	666.93	-1378.2
- citric	1	0.654	667.35	-1377.2
<none>			666.70	-1376.8
- fso2	1	1.829	668.53	-1374.4
- pH	1	4.325	671.02	-1368.5
- tso2	1	8.728	675.43	-1358.0
- chlorides	1	8.761	675.46	-1357.9
- so4	1	27.287	693.98	-1314.7
- vacidity	1	35.000	701.70	-1297.0
- alcohol	1	119.669	786.37	-1114.8

```
Step: AIC=-1378.54
```

```
quality ~ vacidity + citric + rsugar + chlorides + fso2 + tso2 +  
pH + so4 + alcohol
```

	Df	Sum of Sq	RSS	AIC
- rsugar	1	0.257	667.06	-1379.9
- citric	1	0.565	667.37	-1379.2
<none>			666.81	-1378.5
- fso2	1	1.901	668.71	-1376.0

```

- pH          1      7.065 673.87 -1363.7
- chlorides   1      9.940 676.75 -1356.9
- tso2        1     10.031 676.84 -1356.7
- so4         1     27.673 694.48 -1315.5
- vacidity    1     36.234 703.04 -1295.9
- alcohol     1    120.633 787.44 -1114.7

```

Step: AIC=-1379.93

```

quality ~ vacidity + citric + chlorides + fso2 + tso2 + pH +
      so4 + alcohol

```

	Df	Sum of Sq	RSS	AIC
- citric	1	0.475	667.54	-1380.8
<none>			667.06	-1379.9
- fso2	1	2.064	669.13	-1377.0
- pH	1	7.138	674.20	-1364.9
- tso2	1	9.828	676.89	-1358.5
- chlorides	1	9.832	676.89	-1358.5
- so4	1	27.446	694.51	-1317.5
- vacidity	1	35.977	703.04	-1297.9
- alcohol	1	122.667	789.73	-1112.0

Step: AIC=-1380.79

```

quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 + alcohol

```

	Df	Sum of Sq	RSS	AIC
<none>			667.54	-1380.8
- fso2	1	2.394	669.93	-1377.1
- pH	1	7.073	674.61	-1365.9
- tso2	1	10.787	678.32	-1357.2
- chlorides	1	10.809	678.35	-1357.1
- so4	1	27.060	694.60	-1319.2
- vacidity	1	42.318	709.85	-1284.5
- alcohol	1	124.483	792.02	-1109.4

Call:

```

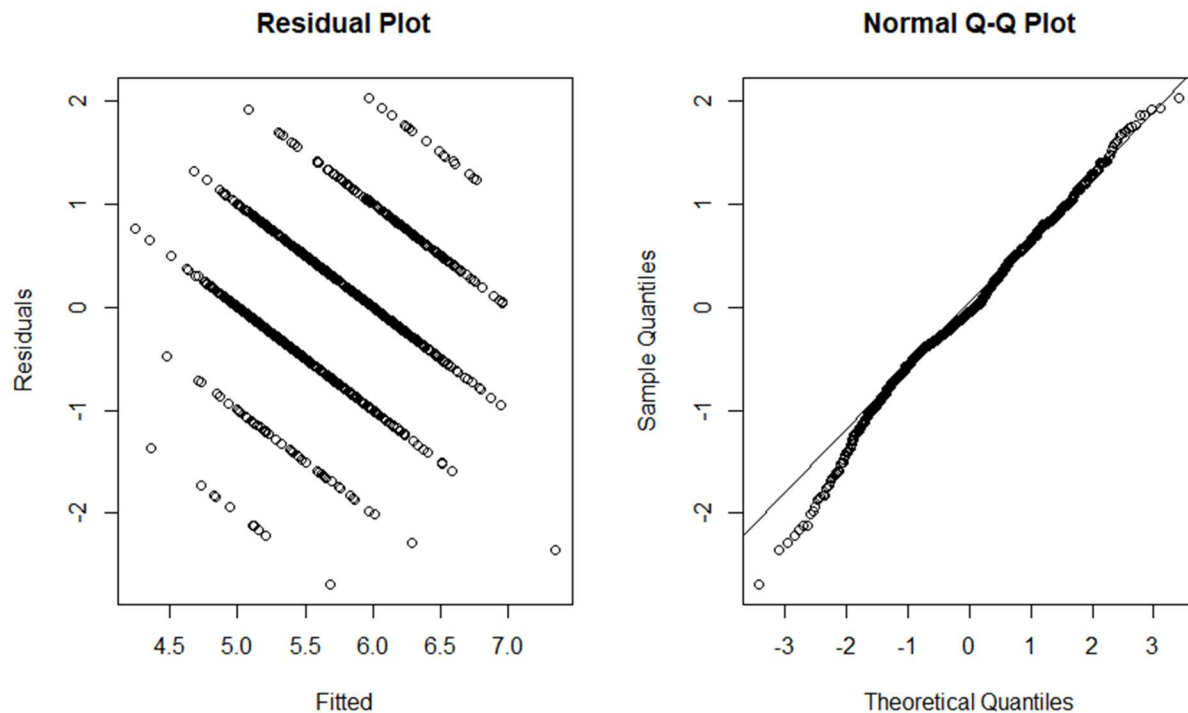
lm(formula = quality ~ vacidity + chlorides + fso2 + tso2 + pH +
      so4 + alcohol)

```

Coefficients:

(Intercept)	vacidity	chlorides	fso2	tso2	
pH	so4	alcohol			
4.430099	-1.012753	-2.017814	0.005077	-0.003482	-
0.482661	0.882665	0.289303			

I check the model assumptions. The two diagnostic plots are presented:



It is difficult to interpret this residual plot. The diagonal streaks arise from discrete values of wine quality. The qqplot suggests that the distribution of residuals is slightly light-tailed on the left tail of distribution.

```
> shapiro.test(residuals(lmod))

      Shapiro-Wilk normality test

data:  residuals(lmod)
W = 0.99137, p-value = 4.321e-08
```

Shapiro test punishes hard for mild deviation from normality and as such gives a verdict that the distribution is clearly not normal.

Recall that for short-tailed distribution the large sample size allows us to make estimation of coefficients, so we have little reason to abandon linear regression. However, it does not allow inference such as confidence intervals.

The model also fails Durbin Watson test implying errors are heavily correlated.


```
> dwtest(quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
alcohol)

      Durbin-Watson test
```

```
data:  quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
alcohol
DW = 1.75, p-value = 2.345e-07
alternative hypothesis: true autocorrelation is greater than 0
As such I use bootstrap confidence intervals for coefficients.
```

One solution is to build covariance structure into the model via GLS, but it is not done here.

Now checking to see if we need to remove any data points, I first make a half-normal plot of leverage. There are two data points (152 and 259) that have high leverages and diverge substantially from the rest of the data.




Next, I compute leave-one-out residuals

```
> tail(jack[order(abs(jack))])
      460      900     1506     1277      653      833
-3.284724 -3.349168 -3.434795 -3.550248 -3.678457 -4.185391
> qt(.05/(1599*2),1599-8)
[1] -4.176048
> jack[152]
      152
-2.851774
```

Considering I used Bonferroni correction which is conservative in finding fewer outliers than the nominal level of confidence would dictate, the data point 833 is way beyond the confidence level.



Next we find the most influential points:



We try excluding this particular point 152, which also had the highest leverage. Recall that 152 had t-score of -2.85 which is quite high.

```
> summary(lm(quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
  alcohol, subset=(cook<max(cook))))
```

Call:

```
lm(formula = quality ~ vacidity + chlorides + fso2 + tso2 + pH +
  so4 + alcohol, subset = (cook < max(cook)))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.71375	-0.36843	-0.04987	0.46154	2.03385

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	4.4002381	0.4021530	10.942	< 2e-16	***
vacidity	-1.0051906	0.1006526	-9.987	< 2e-16	***
chlorides	-1.7757498	0.4056341	-4.378	1.28e-05	***
fso2	0.0053602	0.0021231	2.525	0.0117	*
tso2	-0.0035385	0.0006855	-5.162	2.76e-07	***
pH	-0.4912605	0.1173342	-4.187	2.98e-05	***
so4	0.9136403	0.1101995	8.291	2.37e-16	***

```

alcohol      0.2904650  0.0167632  17.328  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

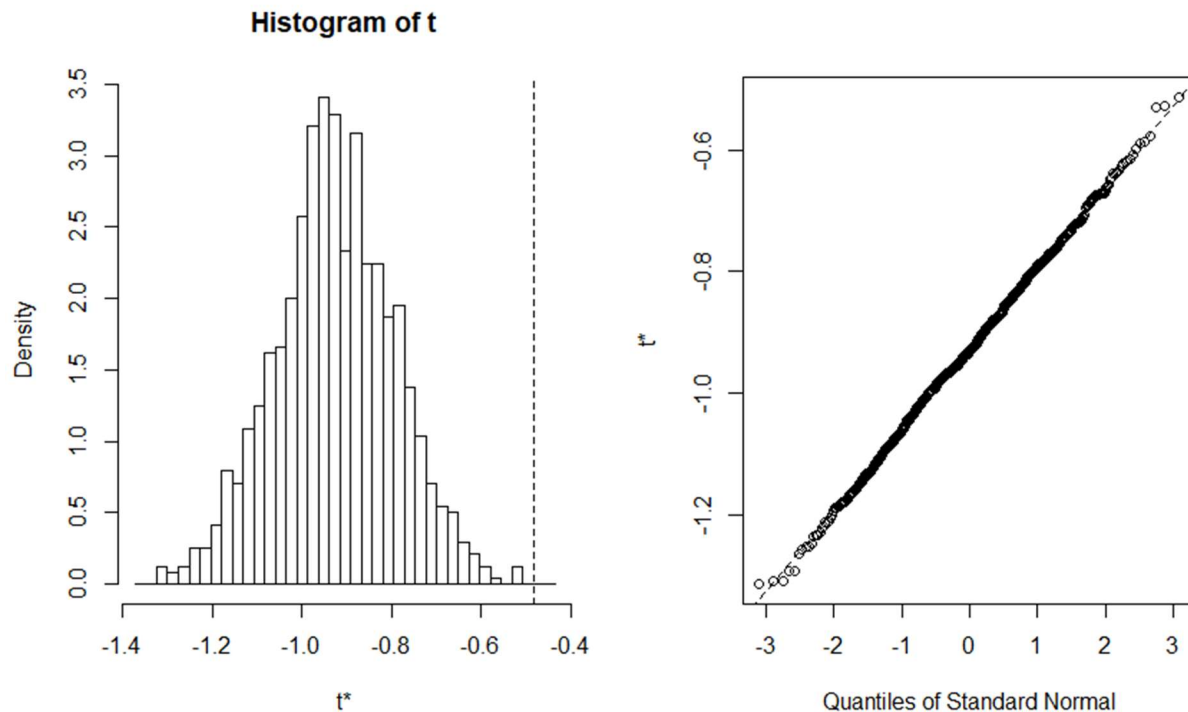
```

Residual standard error: 0.6463 on 1590 degrees of freedom
Multiple R-squared:  0.3611,    Adjusted R-squared:  0.3583
F-statistic: 128.4 on 7 and 1590 DF,  p-value: < 2.2e-16

```

I observed that there is no significant change in the model. P-values are roughly the same for predictors and the estimates do not change. As a final model I decided not to exclude any data point for the reason that they do not change the model substantially and removing outlier automatically without understanding physical context of data can be dangerous.

In order to estimate confidence interval, I use bootstrap method because it does not require any distributional assumptions and provide more accurate inferences when the data are not well behaved.



ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```

boot(data = wine, statistic = bs, R = 1000, formula = quality ~
      vacidity + chlorides + fso2 + tso2 + pH + so4 + alcohol)

```

Bootstrap Statistics :

	original	bias	std. error
t1*	4.430098698	0.433583029	0.4640263454
t2*	-1.012752700	1.005001221	0.1023640512
t3*	-2.017813817	0.912322429	0.4125870793
t4*	0.005077370	-0.004975567	0.0022538235
t5*	-0.003482245	0.003449912	0.0007027508
<b>t6*</b>	<b>-0.482661444</b>	<b>-0.444583503</b>	<b>0.1329182029</b>
t7*	0.882665133	-0.877056724	0.1056710125
t8*	0.289302753	0.088698475	0.0192353832

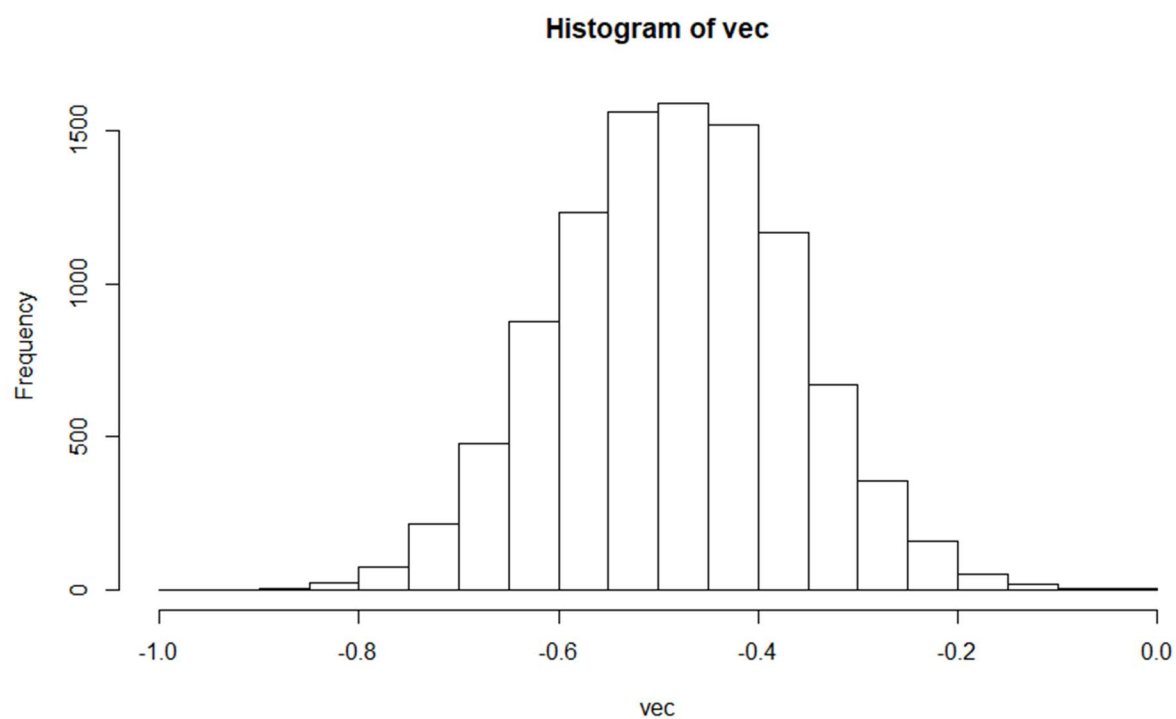
t6 represents the coefficient for pH.

The confidence interval is obtained via:


```
> lmod = lm(quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
alcohol)
> preds = fitted(lmod)
> resids = residuals(lmod)
> vec=numeric(10000)
> for(i in 1:10000) {
+   ynew = preds + sample(resids, rep=TRUE)
+   vec[i]=summary(lm(ynew ~ vacidity + chlorides + fso2 + tso2 + pH +
so4 + alcohol))$coef[6]
+ }
> par(mfrow=c(1,1))

> a = -0.48266
> length(vec[vec>a])/10000
[1] 0.4976
> sort(vec)[250]
[1] -0.7126148
> sort(vec)[9750]
[1] -0.2534445
```

The confidence interval (95%) for the parameter is (-0.7126,-0.2534)




I look to find a suitable box-cox transformation that can change the model. Following is the log-likelihood of the box-coefficient parameter lambda.



The plot suggests that the box-cox parameter is nearly 1. To be more precise,

```
> head(cbind(boxcox(lmod)$x, boxcox(lmod)$y)[order(-boxcox(lmod)$y),])
      [,1]      [,2]
[1,] 0.9090909 -2451.054
[2,] 0.9494949 -2451.103
[3,] 0.8686869 -2451.125
[4,] 0.9898990 -2451.269
[5,] 0.8282828 -2451.315
[6,] 1.0303030 -2451.552
```

The parameter is approximately 0.91.



The 95% confidence interval runs from 0.7 to 1.15. There is no good reason to transform.

However, it might be a good idea to compare the transformed model with the original. With the techniques I learned in class, R-squared and AIC/BIC cannot be used to compare the models that are not nested or use the same set of data. Once the model is transformed by power transform or logarithm, response or predictor variable has different values, and R-squared cannot be compared between a model with untransformed Y and one with transformed Y. Not only that, we know the original model is flawed in that residuals are correlated and slightly non-normal, so we cannot quite compute R-squared.

We are interested in formulating a model that has predictive power. As such, we compare how much less error a model. We can root-mean-squared/ cross-validation approach, where we compare the predictive power of each model.

Here is the box-cox transformed model with the power = 0.91.

Call:

```
lm(formula = quality_ ~ vacidity + chlorides + fso2 + tso2 +
    pH + so4 + alcohol)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.41485	-0.31422	-0.03772	0.40452	1.73805

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	3.395289	0.351431	9.661	< 2e-16	***
vacidity	-0.888548	0.087957	-10.102	< 2e-16	***
chlorides	-1.758634	0.346743	-5.072	4.40e-07	***
fso2	0.004448	0.001854	2.399	0.0165	*
tso2	-0.003014	0.000599	-5.032	5.40e-07	***
pH	-0.418804	0.102536	-4.084	4.64e-05	***
so4	0.766301	0.095864	7.994	2.50e-15	***
alcohol	0.250546	0.014650	17.103	< 2e-16	***

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.565 on 1591 degrees of freedom

Multiple R-squared: 0.3579, Adjusted R-squared: 0.3551

F-statistic: 126.7 on 7 and 1591 DF, p-value: < 2.2e-16

Looking at the diagnostics, the transformation doesn't significantly improve the assumptions.



```
> trainwine <- wine[1:1400,]
> testwine <- wine[1401:1599,]
>
> mod1 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates +
alcohol, data=trainwine)
> mod2 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity +
chlorides + free.sulfur.dioxide + total.sulfur.dioxide + pH +
sulphates + alcohol, data=trainwine)
>
> rmse(predict(mod1,testwine), testwine$quality)
[1] 0.6994192
> rmse(predict(mod2,testwine), testwine$quality)
[1] 1.304175
```

I noticed that the RMSE of the original model is smaller, which means it incurs smaller prediction error. I divided the data into 8 training/testing sets, computed rmse for each, and averaged them. Then I have 0.658 for the first model and 1.367.

```
> (0.699 + 0.631+0.623 + 0.692 + 0.613 + 0.679 +0.666 + 0.662)/8
[1] 0.658125
> (1.304 + 1.269 + 1.495 + 1.412 + 1.296 + 1.425 + 1.53 + 1.207)/8
[1] 1.36725
```

There is a caution to this result however. The root mean squared contains error sigma-squared, and if they are substantially different, it becomes difficult to compare the root mean squared. Looking at the estimated residual standard error, we see that they are slightly different: 0.64 vs 0.56. However, the difference appears small enough that we can compare the root mean squared (data-driven comparison measure) to see which model is better in terms of prediction error given the data.

## II.

To reiterate the problem and identify its components, we have four processes of production from A to D - labeled as “treat”, and the five types of blend are the fixed effect (which will be changed to random effect later)

We run regression of yield on the sum of dummy variables for treatment and blend, which R sets up automatically. Intercept is the reference level (or treat A), and the others are its relative magnitude.

```
> lmod <- lm(yield ~ treat + blend, penicillin)
> summary(lmod)
```

Call:

```
lm(formula = yield ~ treat + blend, data = penicillin)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.00	-2.25	-0.50	2.25	6.00

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	90.000	2.745	32.791	4.1e-13	***
treatB	1.000	2.745	0.364	0.72194	
treatC	5.000	2.745	1.822	0.09351	.
treatD	2.000	2.745	0.729	0.48018	
blendBlend2	-9.000	3.069	-2.933	0.01254	*
blendBlend3	-7.000	3.069	-2.281	0.04159	*
blendBlend4	-4.000	3.069	-1.304	0.21686	
blendBlend5	-10.000	3.069	-3.259	0.00684	**

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.34 on 12 degrees of freedom

Multiple R-squared: 0.5964, Adjusted R-squared: 0.361

F-statistic: 2.534 on 7 and 12 DF, p-value: 0.07535


```
> lmod_ <- lm(yield ~ blend, penicillin)
```

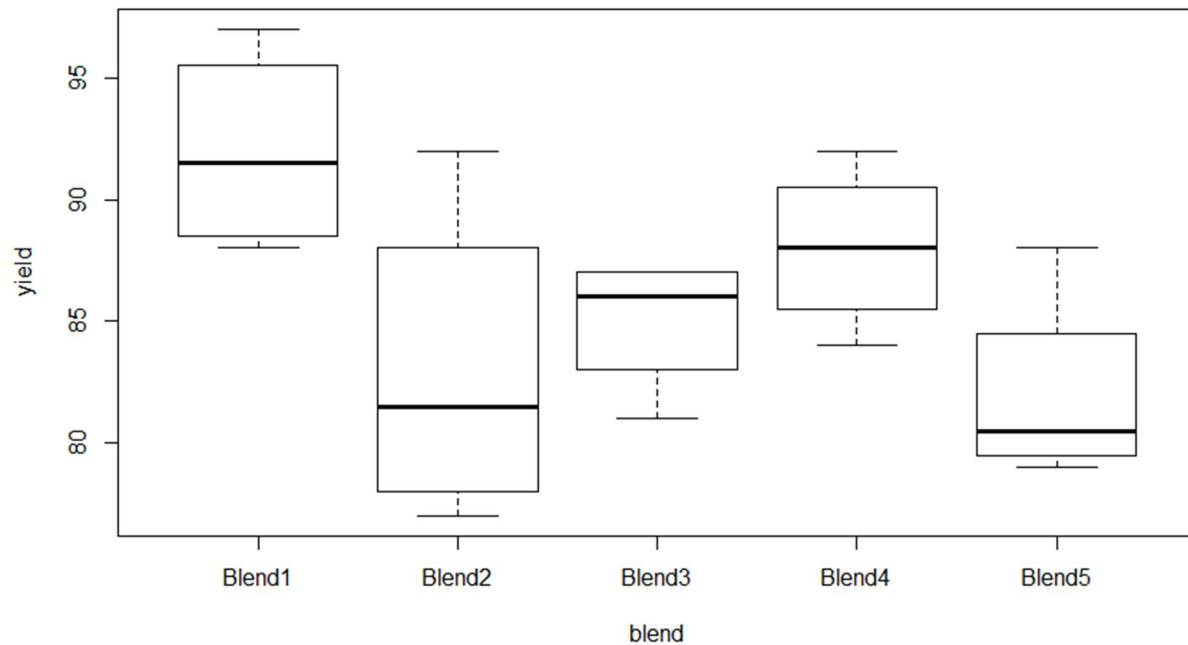


```
> anova(lmod,lmod_)
Analysis of Variance Table

Model 1: yield ~ treat + blend
Model 2: yield ~ blend
      Res.Df RSS Df Sum of Sq      F Pr(>F)
1          12 226
2          15 296 -3        -70 1.2389 0.3387
```

According to the summary of regression, none of the treatments is significant, which means none of the treatments is significantly different from A. When we look at the boxplots of each process with yield, the means are not very different from one another except C, whose p-value is low. On the other hand, there seems to be no block effect except blend 2, 3, and 5, which are significant. As seen in the boxplot, there seem to be significant difference in means among each blend. If we look at the anova test comparing the model without treatment and with treatment, the F-statistic p-value is high, so we can conclude the treatments are not significant.





Using blends increases efficiency as it reduces the variance of the model, where `lmod` is the blocked model whereas the unblocked model is `lmod2` whose variance is underlined. The efficiency is  $30.625/18.833 = 1.62 > 1$ , which implies it is easier to detect treatment effect under RCBD.

```
> anova(lmod)
```

```
Analysis of Variance Table
```

```
Response: yield
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
treat	3	70	23.333	1.2389	0.33866
blend	4	264	66.000	3.5044	0.04075 *
Residuals	12	226	<u>18.833</u>		

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
> anova(lmod2)
```

```
Analysis of Variance Table
```

```
Response: yield
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
treat	3	70	23.333	0.7619	0.5318

Residuals 16      490      30.625

Now I consider blend as the random effect whose noise is correlated with each block, so I am interested in the variance of this random effect. Using `gls` function simply as follows means we are considering both treatment and block as random variables, and is wrong.

```
> gls(yield ~ treat + blend, penicillin)
Generalized least squares fit by REML
Model: yield ~ treat + blend
Data: penicillin
Log-restricted-likelihood: -39.82778
```

```
Coefficients:
(Intercept)      treatB      treatC      treatD blendBlend2
blendBlend3 blendBlend4 blendBlend5
              90          1          5          2          -9          -
7          -4          -10
```

```
Degrees of freedom: 20 total; 12 residual
Residual standard error: 4.339739
```

If we consider RCBD with one random effect and the variance of random effect, we can regress yield on treat with `gls` and treat its residual error as the sum of sigma-squared (variance of OLS) and v-squared (variance of random effect). Thus, we compute the residual variance of GLS and subtract by sigma-squared (variance of OLS), which is v-squared (variance of random effect). We have  $v^2 = 5.53^2 - 4.34^2 = 12$ .

```
> glmod <- gls(yield ~ treat, penicillin)
> summary(glmod) #5.533 = 30.61
Generalized least squares fit by REML
Model: yield ~ treat
Data: penicillin
      AIC      BIC    logLik
116.5929 120.4558 -53.29643
```

```
Coefficients:
      Value Std.Error t-value p-value
(Intercept)   84  2.474874 33.94113  0.0000
treatB         1  3.500000  0.28571  0.7788
treatC         5  3.500000  1.42857  0.1724
treatD         2  3.500000  0.57143  0.5756
```

```
Correlation:
      (Intr) treatB treatC
treatB -0.707
```

```
treatC -0.707  0.500
treatD -0.707  0.500  0.500
```

Standardized residuals:

	Min	Q1	Med	Q3	Max
	-1.6263142	-0.5872801	0.0000000	0.5421047	1.4456126

Residual standard error: **5.533986**

Degrees of freedom: 20 total; 16 residual

```
> lmod <- lm(yield ~ treat + blend, penicillin)
> summary(lmod) #4.34 = 18.83
```

Call:

```
lm(formula = yield ~ treat + blend, data = penicillin)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-5.00	-2.25	-0.50	2.25	6.00

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	90.000	2.745	32.791	4.1e-13	***
treatB	1.000	2.745	0.364	0.72194	
treatC	5.000	2.745	1.822	0.09351	.
treatD	2.000	2.745	0.729	0.48018	
blendBlend2	-9.000	3.069	-2.933	0.01254	*
blendBlend3	-7.000	3.069	-2.281	0.04159	*
blendBlend4	-4.000	3.069	-1.304	0.21686	
blendBlend5	-10.000	3.069	-3.259	0.00684	**

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: **4.34** on 12 degrees of freedom

Multiple R-squared: 0.5964, Adjusted R-squared: 0.361

F-statistic: 2.534 on 7 and 12 DF, p-value: 0.07535

### III.

The idea of transformation is attached. I implement the idea where I re-defined  $y_*$  as  $y - x^2$  and  $x_*$  as  $x1 - x2$ , and obtain the coefficients as follows:

Call:

```
lm(formula = y_ ~ x_)
```

Coefficients:

(Intercept)	<u>3.1956</u>	<u>0.3668</u>
		$x_{-}$

Call:

lm(formula =  $y_{-} \sim x_{-}$ )

Residuals:

Min	1Q	Median	3Q	Max
-2.93143	-0.78183	0.01373	0.86331	2.49095

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	3.19563	0.14641	21.827	< 2e-16 ***
$x_{-}$	0.36676	0.09119	4.022	0.000114 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.1 on 98 degrees of freedom

Multiple R-squared: 0.1417, Adjusted R-squared: 0.1329

F-statistic: 16.18 on 1 and 98 DF, p-value: 0.0001136

#### IV

#### R Code

```
# 1
load("C:/Users/jihun/Downloads/winequality.RData")
quality <- wine$quality
facidity <- wine$fixed.acidity
vacidity <- wine$volatile.acidity
citric <- wine$citric.acid
rsugar <- wine$residual.sugar
chlorides <- wine$chlorides
fso2 <- wine$free.sulfur.dioxide
tso2 <- wine$total.sulfur.dioxide
density <- wine$density
pH <- wine$pH
so4 <- wine$sulphates
alcohol <- wine$alcohol
# linear model
```

```

lmod <- lm(quality ~ facidity + vacidity + citric + rsugar + chlorides
+ fso2 + tso2 + density + pH + so4 + alcohol)
summary(lmod)
# model selection
step(lmod)
lmod4 <- lm(quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
alcohol)
summary(lmod4)
lmod=lmod4
# diagnostics
# homoscedasticity
par(mfrow=c(1,2))
plot(fitted(lmod),residuals(lmod),xlab="Fitted",ylab="Residuals",main=
"Residual Plot")
qqnorm(residuals(lmod))
qqline(residuals(lmod))
# normality
shapiro.test(residuals(lmod))
# independence
library(lmtest)
dwtest(quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
alcohol)

# outliers
# check leverages
halfnorm(lm.influence(lmod)$hat,ylab="Leverages")
# outlier
jack <- rstudent(lmod)
jack[which.max(abs(jack))]
tail(jack[order(abs(jack))])
qt(.05/(1599*2),1599-8)
jack[152]
# influence
cook <- cooks.distance(lmod)
halfnorm(cook, ylab="Cook's distances")
summary(lm(quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
alcohol, subset=(cook<max(cook))))

# bootstrap
library(boot)
# function to obtain regression weights
bs <- function(formula, data, indices) {
  d <- data[indices,] # allows boot to select sample

```

```

    fit <- lm(formula, data=d)
    return(coef(fit))
}
# bootstrapping with 1000 replications
results <- boot(data=wine, statistic=bs,
                R=1000, formula=quality ~ vacidity + chlorides + fso2
+ tso2 + pH + so4 + alcohol)

# view results
results
plot(results, index=6) # pH

```

```

lmod = lm(quality ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
alcohol)
preds = fitted(lmod)
resids = residuals(lmod)
vec=numeric(10000)
for(i in 1:10000) {
  ynew = preds + sample(resids, rep=TRUE)
  vec[i]=summary(lm(ynew ~ vacidity + chlorides + fso2 + tso2 + pH +
so4 + alcohol))$coef[6]
}
par(mfrow=c(1,1))
hist(vec)
a = -0.48266
length(vec[vec>a])/10000
sort(vec)[250]
sort(vec)[9750]
# Box-cox transformation
library(MASS)
boxcox(lmod, plotit=T)
head(cbind(boxcox(lmod)$x, boxcox(lmod)$y)[order(-boxcox(lmod)$y),])
boxcox(lmod,plotit=T,lambda=seq(0.5,1.5,by=0.1))

# data-driven approach

```

```

quality_ = ((quality -1)^0.91)/0.91
fit_ = lm(quality_ ~ vacidity + chlorides + fso2 + tso2 + pH + so4 +
alcohol)
summary(fit_)
par(mfrow=c(1,2))
plot(fit_ $fitted.values, fit_ $residuals, main="Transformed Fitted Value
vs Residuals")

```

```

qqnorm(fit_$residuals)
qqline(residuals(fit_))

library(Metrics)

trainwine <- wine[1:1400,]
testwine <- wine[1401:1599,]

mod1 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine)
mod2 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine)

rmse(predict(mod1,testwine), testwine$quality) # 0.6524137
rmse(predict(mod2,testwine), testwine$quality) # 1.283869

trainwine2 <- wine[-c(1201:1400),]
testwine2 <- wine[c(1201:1400),]
mod21 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine2)
mod22 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine2)

rmse(predict(mod21,testwine2), testwine2$quality) # 0.631
rmse(predict(mod22,testwine2), testwine2$quality) # 1.269

trainwine3 <- wine[-c(1001:1200),]
testwine3 <- wine[c(1001:1200),]
mod31 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine3)
mod32 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine3)

rmse(predict(mod31,testwine3), testwine3$quality) # 0.623
rmse(predict(mod32,testwine3), testwine3$quality) # 1.495

trainwine4 <- wine[-c(801:1000),]
testwine4 <- wine[c(801:1000),]

```



```

mod41 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine4)
mod42 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine4)

rmse(predict(mod41,testwine4), testwine4$quality) # 0.692
rmse(predict(mod42,testwine4), testwine4$quality) # 1.412

trainwine5 <- wine[-c(601:800),]
testwine5 <- wine[c(601:800),]
mod51 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine5)
mod52 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine5)

rmse(predict(mod51,testwine5), testwine5$quality) # 0.613
rmse(predict(mod52,testwine5), testwine5$quality) #1.296

trainwine6 <- wine[-c(401:600),]
testwine6 <- wine[c(401:600),]
mod61 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine6)
mod62 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine6)

rmse(predict(mod61,testwine6), testwine6$quality) # 0.679
rmse(predict(mod62,testwine6), testwine6$quality) # 1.425

trainwine7 <- wine[-c(201:400),]
testwine7 <- wine[c(201:400),]
mod71 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine7)
mod72 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine7)

rmse(predict(mod71,testwine7), testwine7$quality) # 0.666
rmse(predict(mod72,testwine7), testwine7$quality) # 1.53

```

```

trainwine8 <- wine[-c(1:200),]
testwine8 <- wine[c(1:200),]
mod81 <- lm(quality ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine8)
mod82 <- lm(((quality -1)^0.91)/0.91 ~ volatile.acidity + chlorides +
free.sulfur.dioxide + total.sulfur.dioxide + pH + sulphates + alcohol,
data=trainwine8)

rmse(predict(mod81,testwine8), testwine8$quality) # 0.662
rmse(predict(mod82,testwine8), testwine8$quality) # 1.207

library(analogue)

pcrmod <- pcr(quality ~ volatile.acidity + chlorides +
              free.sulfur.dioxide + total.sulfur.dioxide +
              pH + sulphates + alcohol, data=trainwine,
              validation="CV",ncomp=50)

rmse(predict(pcrmod, testwine$quality))

# 2
library(faraway)
data(penicillin)
penicillin
plot(yield ~ treat, penicillin,pch=unclass(blend))
plot(yield ~ blend, penicillin,pch=unclass(treat))
# RCBD
lmod <- lm(yield ~ treat + blend, penicillin)
summary(lmod)
lmod_ <- lm(yield ~ blend, penicillin)
anova(lmod,lmod_)

# CRD
lmod2 <- lm(yield ~ treat, penicillin)
summary(lmod2)
anova(lmod2)

# random effect

library(nlme)

```

```
gls(yield ~ treat + blend, penicillin)

glmod <- gls(yield ~ treat, penicillin)
summary(glmod) #5.533 = 30.61

lmod <- lm(yield ~ treat + blend, penicillin)
summary(lmod) #4.34 = 18.83

# 3
load("C:/Users/jihun/Downloads/Constrained.RData")
y_ = y - x2
x_ = x1 - x2
summary(lm(y_ ~ x_))
```